

Supervised Learning Exam

Università degli Studi di Milano-Bicocca

12-06-2023

Monaco Filippo – 840089
Picione Marco - 827116

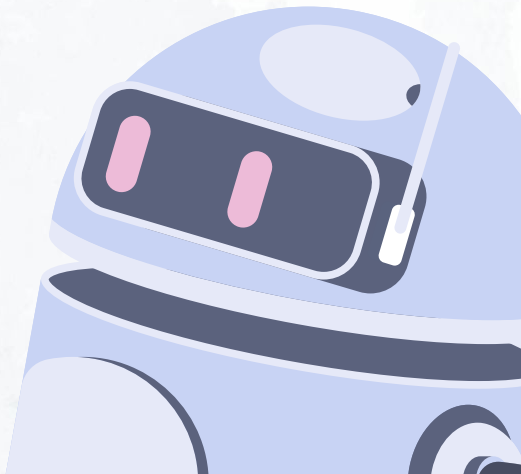


Table of contents

- 00 → Introduction
- 01 → Dataset
- 02 → Bag of Visual Words
- 03 → Convolutional Neural Network
- 04 → Conclusion

00 →

Introduction

AIM: Multi-class image classification
MODELS: Visual Bag of Words & Convolutional Neural Network



Multi-class Image Classification

Computer vision task that aims to assign a single label to an image.



dog?

ant?

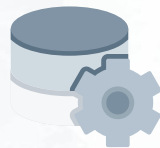
cat.

chair?

car?

01 →

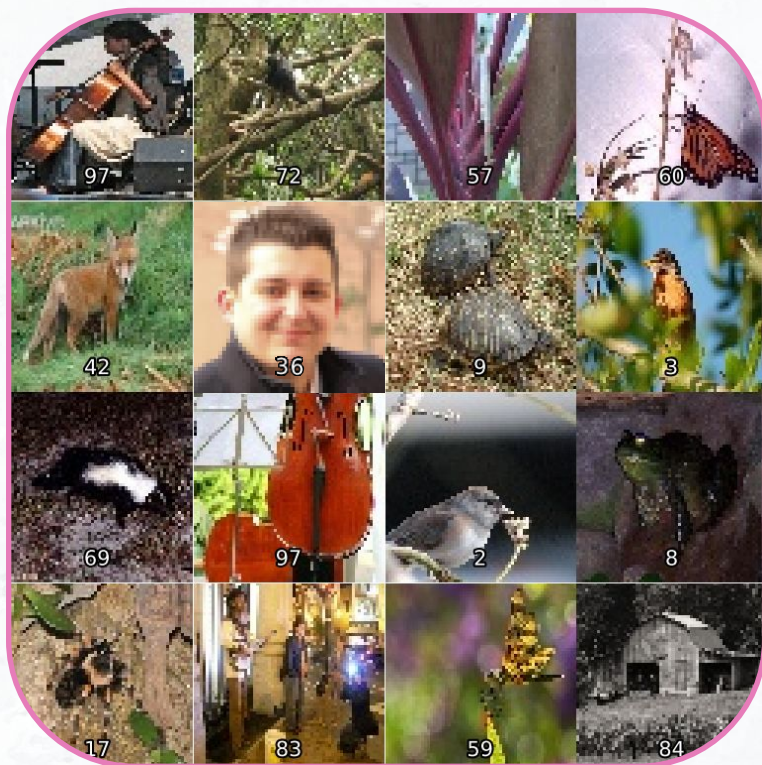
Dataset



TinyImageNet, a reduced version of the widely-used ImageNet

- 100 classes
- 1000 images per class
- 64 x 64 image size

Data was split 80% for **Training**, 20% for **Validation**. A different set for **Testing**.



Some images of the dataset with associated label.

Note that the low resolution is due to the low size of the images (64x64).

Data pre-processing

Bag of Words:

- RGB to **Gray**Scale

Convolutional Neural Network:

- **Resize** to 256 x 256
- Central **Crop** to 224 x 224
- Image to **Tensor**
- **Normalization** (mean [0.485, 0.456, 0.406] , std [0.229, 0.224, 0.225])
- **Batching** (64 images per batch)

Data Augmentation

Certain classes performed much worse than others.

Our Neural Network particularly did not enjoy pictures of cockroaches.

Augment this class by creating new images with random transformations:

- random Horizontal/Vertical **Flip**
- random **Brightness** variation
- random **Rotation**



ugly.

02 →

Bag of Words



Visual Bag of Words tries to represent images by a set of features (**keypoints** and **descriptors**), which are used to construct a visual **vocabulary**.

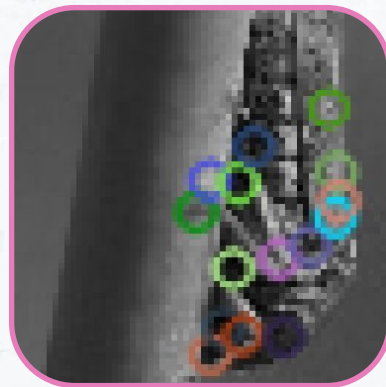
The image is described by a frequency **histogram** of features, which is then used for classification.

Extracting Local Features

SIFT is used to extract local features from images:

- **keypoints** (stand-out points)
- **descriptors** (description of keypoints)

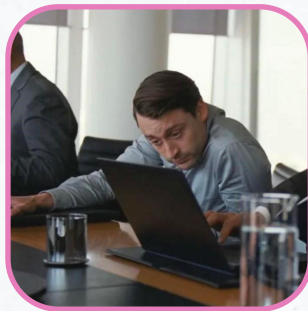
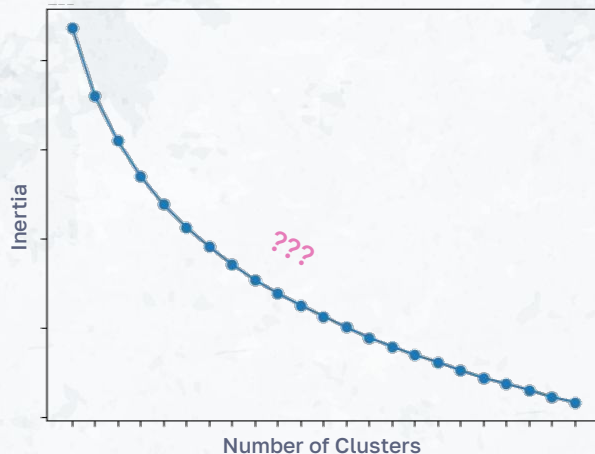
37 keypoints per image (on average), each described by a 128-dimensional vector.



Keypoints Clustering

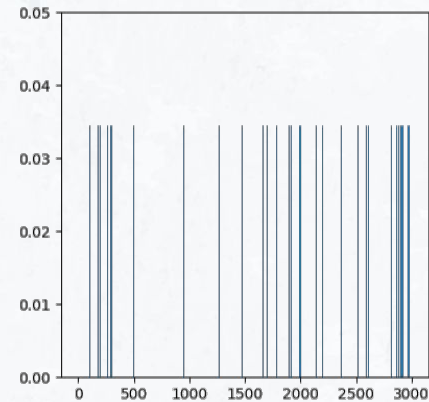
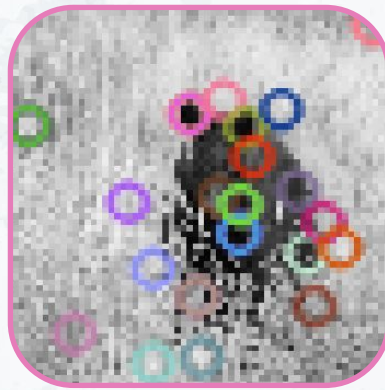
Clustering of all descriptors using an unsupervised algorithm:
MiniBatches k-Means.

Elbow point method was inconclusive:
arbitrarily chosen **3000** clusters (number of visual words).



Visual Histogram

Descriptors are clustered to construct a Visual **Histogram** based on the frequency of features w.r.t. the visual vocabulary.



Classifier

Histograms are used to train a traditional **classifier**, which is then used to predict labels on the Test set.

Classifiers tested:

- k-Nearest Classifier
- Stochastic Gradient Descent Classifier
- Support Vector Classifier
- Decision Tree
- Random Forest
- AdaBoost
- Bagging

	SVC	SGDC
Accuracy	6.7 %	4.7 %
Training Time	4 hours	4 minutes
Testing Time	1 hour	0.1 seconds

Motivation

Very **poor** performance achieved.

Possible reasons:

- The visual vocabulary is not optimal
- **Overfitting**

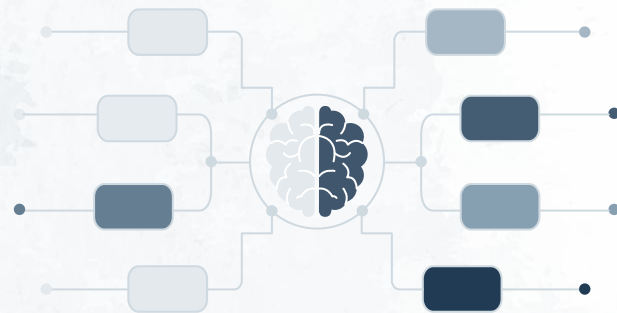
Further testing would be required.



03 →

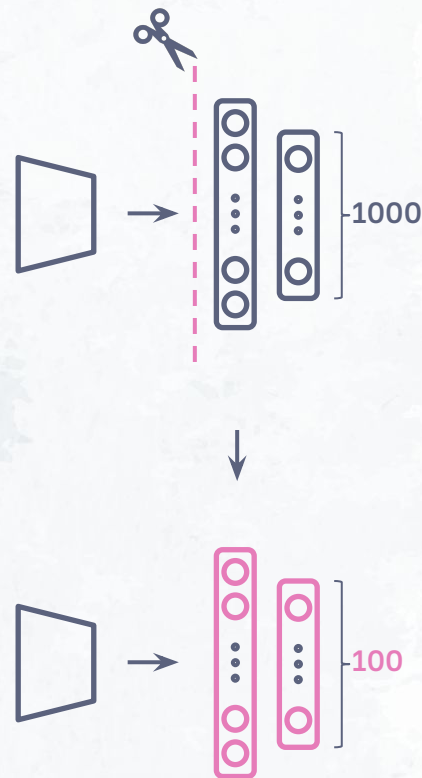
Convolutional Neural Network

Biologically-inspired Networks of neurons widely used in computer vision tasks. Stack multiple feature extractors as hidden neural layers and train them using ground truth labels.



Transfer Learning

1. Take **pre-trained** model
 - **efficientnet_b0**, trained on ImageNet
2. Change output **classifier**
 - from 1000 to **100** classes
3. **Freeze** main model
 - train **only** the newly-added classifier
4. Data **pre-processing**
 - **Resize** to 256 x 256
 - Central **Crop** to 224 x 224
 - Image to **Tensor**
 - **Normalization** (mean [0.485, 0.456, 0.406], std [0.229, 0.224, 0.225])
 - **Batching** (64 images per batch)

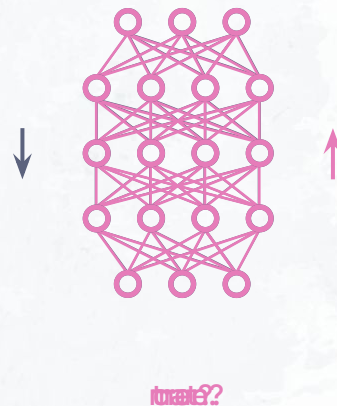


Training and Validation

1. **Input** images to network with corresponding label
2. **Output** logits
3. Compute Cross-Entropy **Loss**
4. **Backpropagate** error, updating classifier parameters
5. **Validate** model at current epoch
6. **Repeat** process until validation loss is stable
7. Save **best** model

Hyper-parameters:

- Epochs : **10**
- Patience : **3**
- Optimizer : **Adam**
- Learning Rate : **0.001**
- Scheduler : **CosineAnnealingLR**

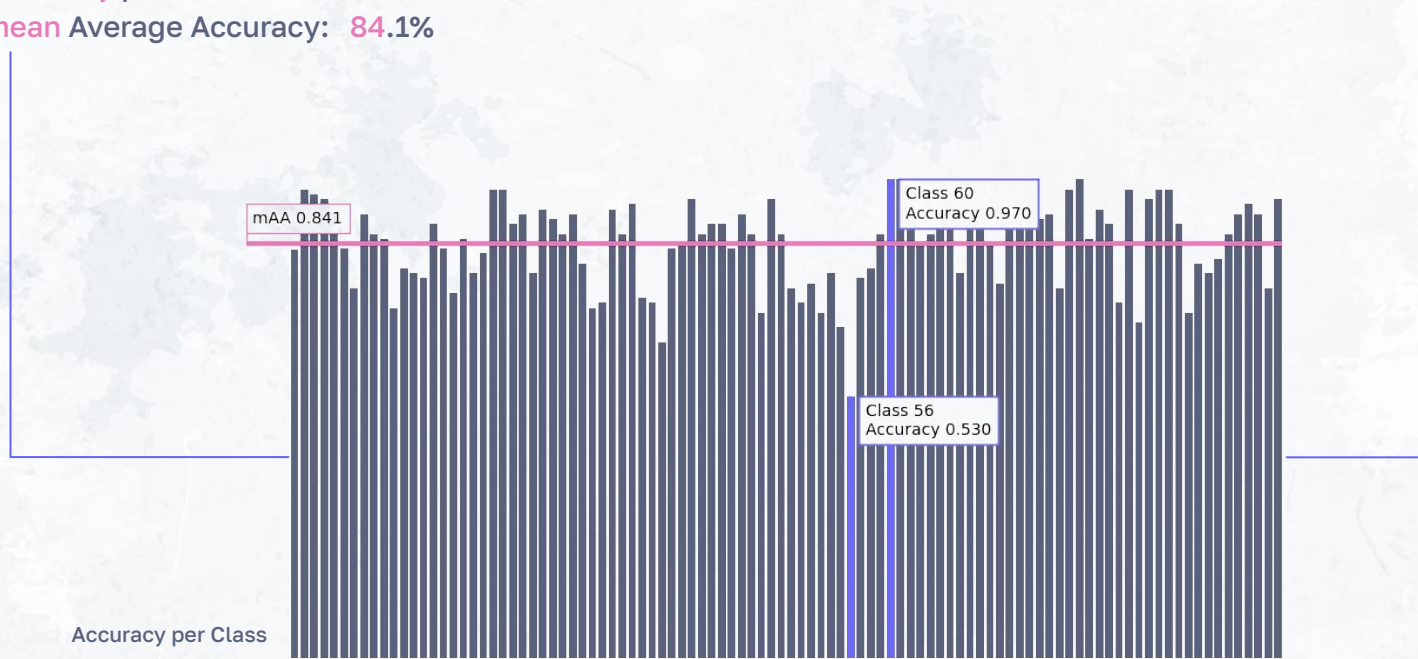


Testing

Metrics:

- **Accuracy** per class
- **mean** Average Accuracy: **84.1%**

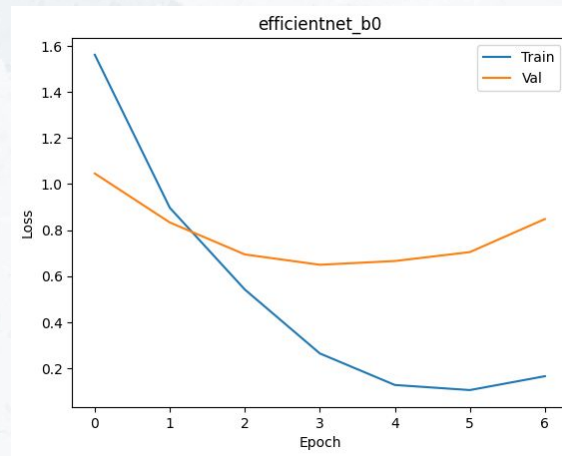
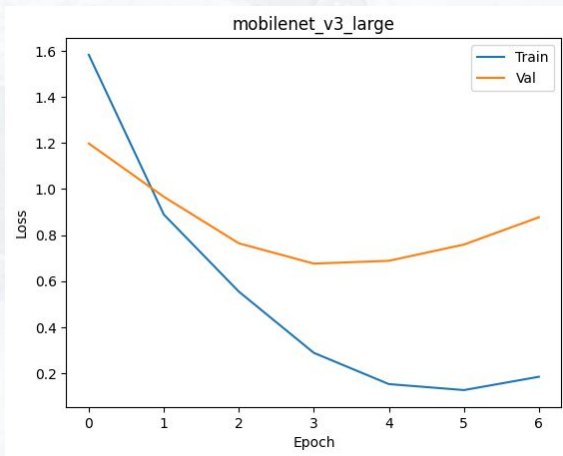
Data augmentation was successful on **class 56**:
from **53.0%** to **75.0%** accuracy.



Hyperparameter Tuning

To achieve the best possible performance, we tried many different pre-trained models:

- **mobilenet_v3_large** 79% accuracy
- **efficientnet_b0** 84% accuracy
- **vgg16_bn** 70% accuracy



Optuna



Python framework for **automatic hyperparameter optimization**.

Sample different sets of values for each parameter and train the network, try to minimize the loss in validation.

Hyperparameters:

- Number of Hidden **Neurons** in Classifier (first and second hidden layer)
- **Dropout** chance
- Training **Epochs**
- **Learning Rate**
- **Optimizer**



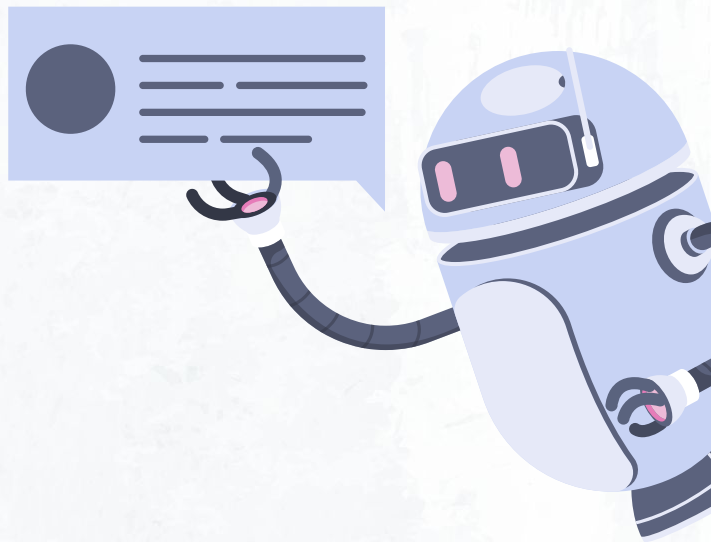
Colab did not want us to reach the maximum of our abilities.

04 →

Conclusion

The Convolutional Neural network works much much better than the traditional classifier, both in terms of overall accuracy, and training times.

And them's the facts.



Thanks!

If you have any more questions, please direct them to

chat.openai.com

(We did not make use of ChatGPT or any other Natural Language Processing models for this project).



Credit:

Filippo Monaco & Marco Picione

under the wise supervision of Prof. Simone Bianco & Mirko Agarla